



Reshaping Data

In this module, we will show you how to:

1. Reshaping data from wide (fat) to long (tall)
2. Reshaping data from long (tall) to wide (fat)
3. Merging Data
4. Perform operations by a grouping variable

What is wide/long data?

See http://www.cookbook-r.com/Manipulating_data/Converting_data_between_wide_and_long_format/

- Wide - multiple columns per observation
 - e.g. visit1, visit2, visit3

```
# A tibble: 2 x 4
  id visit1 visit2 visit3
<int> <dbl> <dbl> <dbl>
1     1     10      4      3
2     2      5      6     NA
```

- Long - multiple rows per observation

```
# A tibble: 5 x 3
  id visit value
<dbl> <int> <dbl>
1     1     1    10
2     1     2     4
3     1     3     3
4     2     1     5
5     2     2     6
```

What is wide/long data?

More accurately, data is wide or long **with respect** to certain variables.

Data used: Charm City Circulator

http://johnmuschelli.com/intro_to_r/data/Charm_City_Circulator_Ridership.csv

```
circ = read_csv(  
  paste0("http://johnmuschelli.com/intro_to_r/",  
        "data/Charm_City_Circulator_Ridership.csv"))  
head(circ, 2)
```

```
# A tibble: 2 x 15  
  day   date orangeBoardings orangeAlightings orangeAverage purpleBoardings  
  <chr> <chr>          <dbl>          <dbl>          <dbl>          <dbl>  
1 Mond... 01/1...          877          1027          952           NA  
2 Tues... 01/1...          777          815          796           NA  
# ... with 9 more variables: purpleAlightings <dbl>, purpleAverage <dbl>,  
#   greenBoardings <dbl>, greenAlightings <dbl>, greenAverage <dbl>,  
#   bannerBoardings <dbl>, bannerAlightings <dbl>, bannerAverage <dbl>,  
#   daily <dbl>
```

```
class(circ$date)
```

```
[1] "character"
```

Creating a Date class from a character date

```
library(lubridate) # great for dates!
```

```
sum(is.na(circ$date))
```

```
[1] 0
```

```
sum( circ$date == "")
```

```
[1] 0
```

```
circ = mutate(circ, date = mdy(date))  
sum( is.na(circ$date) ) # all converted correctly
```

```
[1] 0
```

```
head(circ$date, 3)
```

```
[1] "2010-01-11" "2010-01-12" "2010-01-13"
```

```
class(circ$date)
```

```
[1] "Date"
```

Reshaping data from wide (fat) to long (tall): base R

The `reshape` command exists. It is a **confusing** function. Don't use it.

tidyr package

`tidyr` allows you to “tidy” your data. We will be talking about:

- `gather` - make multiple columns into variables, (wide to long)
- `spread` - make a variable into multiple columns, (long to wide)
- `separate` - string into multiple columns
- `unite` - multiple columns into one string

Reshaping data from wide (fat) to long (tall): tidyr

`tidyr::gather` - puts column data into rows.

We want the column names into “var” variable in the output dataset and the value in “number” variable. We then describe which columns we want to “gather:”

```
long = gather(circ, key = "var", value = "number",  
              -day, -date, -daily)  
head(long, 4)
```

```
# A tibble: 4 x 5  
  day      date      daily var      number  
  <chr>   <date>    <dbl> <chr>    <dbl>  
1 Monday 2010-01-11  952 orangeBoardings 877  
2 Tuesday 2010-01-12  796 orangeBoardings 777  
3 Wednesday 2010-01-13 1212. orangeBoardings 1203  
4 Thursday 2010-01-14 1214. orangeBoardings 1194
```

Reshaping data from wide (fat) to long (tall): tidyr

- Could be explicit on what we want to gather

```
long = gather(circ, key = "var", value = "number",
              starts_with("orange"), starts_with("purple"),
              starts_with("green"), starts_with("banner"))
long
```

```
# A tibble: 13,752 x 5
  day      date      daily var      number
<chr>   <date>   <dbl> <chr>   <dbl>
1 Monday 2010-01-11   952 orangeBoardings 877
2 Tuesday 2010-01-12   796 orangeBoardings 777
3 Wednesday 2010-01-13 1212. orangeBoardings 1203
4 Thursday 2010-01-14 1214. orangeBoardings 1194
5 Friday 2010-01-15 1644. orangeBoardings 1645
6 Saturday 2010-01-16 1490. orangeBoardings 1457
7 Sunday 2010-01-17  888. orangeBoardings 839
8 Monday 2010-01-18  999. orangeBoardings 999
9 Tuesday 2010-01-19 1035. orangeBoardings 1023
10 Wednesday 2010-01-20 1396. orangeBoardings 1375
# ... with 13,742 more rows
```

Reshaping data from wide (fat) to long (tall): tidyr

```
long %>% count(var)
```

```
# A tibble: 12 x 2
  var                n
  <chr>             <int>
1 bannerAlightings  1146
2 bannerAverage     1146
3 bannerBoardings   1146
4 greenAlightings   1146
5 greenAverage      1146
6 greenBoardings    1146
7 orangeAlightings  1146
8 orangeAverage     1146
9 orangeBoardings   1146
10 purpleAlightings 1146
11 purpleAverage     1146
12 purpleBoardings  1146
```

Lab Part 1

[Website](#)

Making a separator

We will use `str_replace` from `stringr` to put `_` in the names

```
long = long %>% mutate(  
  var = var %>%  
    str_replace("Board", " _Board") %>%  
    str_replace("Alight", " _Alight") %>%  
    str_replace("Average", " _Average")  
)  
long %>% count(var)
```

```
# A tibble: 12 x 2  
  var          n  
  <chr>      <int>  
1 banner_Alightings 1146  
2 banner_Average    1146  
3 banner_Boardings  1146  
4 green_Alightings  1146  
5 green_Average     1146  
6 green_Boardings   1146  
7 orange_Alightings 1146  
8 orange_Average    1146  
9 orange_Boardings  1146  
10 purple_Alightings 1146  
11 purple_Average    1146  
12 purple_Boardings  1146
```

Reshaping data from wide (fat) to long (tall): tidyr

Now each `var` is boardings, averages, or alightings. We want to separate these so we can have these by line. Remember `"."` is special character:

```
long = separate(long, var, into = c("line", "type"), sep = "_")
head(long, 2)
```

```
# A tibble: 2 x 6
  day      date      daily line  type      number
<chr> <date>    <dbl> <chr> <chr>    <dbl>
1 Monday 2010-01-11    952 orange Boardings    877
2 Tuesday 2010-01-12    796 orange Boardings    777
```

```
unique(long$line)
```

```
[1] "orange" "purple" "green"  "banner"
```

```
unique(long$type)
```

```
[1] "Boardings" "Alightings" "Average"
```

Re-uniting all the lines

If we had the opposite problem, we could use the `unite` function:

```
reunited = long %>%  
  unite(col = var, line, type, sep = "_")  
reunited %>% select(day, var) %>% head(3) %>% print
```

```
# A tibble: 3 x 2  
  day      var  
  <chr>    <chr>  
1 Monday  orange_Boardings  
2 Tuesday orange_Boardings  
3 Wednesday orange_Boardings
```

We could also use `paste/paste0`.

Lab Part 2

[Website](#)

Reshaping data from long (tall) to wide (fat): tidyr

In `tidyr`, the `spread` function spreads rows into columns. Now we have a long data set, but we want to separate the Average, Alightings and Boardings into different columns:

```
# have to remove missing days
wide = long %>% filter(!is.na(date))
wide = wide %>% spread(type, number)
head(wide)
```

```
# A tibble: 6 x 7
  day   date      daily line Alightings Average Boardings
  <chr> <date>      <dbl> <chr>      <dbl>      <dbl>      <dbl>
1 Friday 2010-01-15 1644 banner         NA         NA         NA
2 Friday 2010-01-15 1644 green         NA         NA         NA
3 Friday 2010-01-15 1644 orange    1643    1644    1645
4 Friday 2010-01-15 1644 purple         NA         NA         NA
5 Friday 2010-01-22 1394 banner         NA         NA         NA
6 Friday 2010-01-22 1394 green         NA         NA         NA
```

Lab Part 3

[Website](#)

Merging: Simple Data

base has baseline data for ids 1 to 10 and Age

```
base <- tibble(id = 1:10, Age = seq(55, 60, length=10))  
head(base, 2)
```

```
# A tibble: 2 x 2  
  id   Age  
<int> <dbl>  
1     1  55  
2     2 55.6
```

visits has ids 1 to 8, then 11 (new id), and 3 visits and outcome

```
visits <- tibble(id = c(rep(1:8, 3), 11), visit= c(rep(1:3, 8), 3),  
                 Outcome = seq(10, 50, length=25))  
tail(visits, 2)
```

```
# A tibble: 2 x 3  
  id visit Outcome  
<dbl> <dbl>   <dbl>  
1     8     3    48.3  
2    11     3    50
```

Joining in `dplyr`

- Merging/joining data sets together - usually on key variables, usually “id”
- `?join` - see different types of joining for `dplyr`
- Let's look at <https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>
- `inner_join(x, y)` - only rows that match for `x` and `y` are kept
- `full_join(x, y)` - all rows of `x` and `y` are kept
- `left_join(x, y)` - all rows of `x` are kept even if not merged with `y`
- `right_join(x, y)` - all rows of `y` are kept even if not merged with `x`

Inner Join

```
ij = inner_join(base, visits)
```

```
Joining, by = "id"
```

```
dim(ij)
```

```
[1] 24  4
```

```
tail(ij)
```

```
# A tibble: 6 x 4
   id    Age visit Outcome
  <dbl> <dbl> <dbl>   <dbl>
1     7  58.3     1     20
2     7  58.3     3    33.3
3     7  58.3     2    46.7
4     8  58.9     2    21.7
5     8  58.9     1     35
6     8  58.9     3    48.3
```

Left Join

```
lj = left_join(base, visits)
```

```
Joining, by = "id"
```

```
dim(lj)
```

```
[1] 26  4
```

```
tail(lj)
```

```
# A tibble: 6 x 4
   id   Age visit Outcome
  <dbl> <dbl> <dbl>   <dbl>
1     7  58.3     2    46.7
2     8  58.9     2    21.7
3     8  58.9     1     35
4     8  58.9     3    48.3
5     9  59.4    NA     NA
6    10  60     NA     NA
```

Logging the joins

The `tidylog` package can show you log outputs from `dplyr` (newly added). You will need to install to use.

```
library(tidylog)
left_join(base, visits)
```

```
Joining, by = "id"left_join: added 2 columns (visit, Outcome)
  > rows only in x      2
  > rows only in y    ( 1)
  > matched rows      24    (includes duplicates)
  >
  > rows total         26
```

```
# A tibble: 26 x 4
   id   Age visit Outcome
<dbl> <dbl> <dbl>   <dbl>
1     1   55     1      10
2     1   55     3     23.3
3     1   55     2     36.7
4     2  55.6     2     11.7
5     2  55.6     1      25
6     2  55.6     3     38.3
7     3  56.1     3     13.3
8     3  56.1     2     26.7
9     3  56.1     1      40
10    4  56.7     1      15
# ... with 16 more rows
```


Right Join

```
rj = right_join(base, visits)
```

```
Joining, by = "id"right_join: added 2 columns (visit, Outcome)
> rows only in x    ( 2)
> rows only in y      1
> matched rows      24
>                    ====
> rows total         25
```

```
dim(rj)
```

```
[1] 25  4
```

```
tail(rj)
```

```
# A tibble: 6 x 4
   id   Age visit Outcome
<dbl> <dbl> <dbl>   <dbl>
1     4  56.7     2    41.7
2     5  57.2     3    43.3
3     6  57.8     1     45
4     7  58.3     2    46.7
5     8  58.9     3    48.3
6    11   NA     3     50
```

Right Join: Switching arguments

```
rj2 = right_join(visits, base)
```

```
Joining, by = "id"right_join: added one column (Age)
> rows only in x   ( 1)
> rows only in y    2
> matched rows     24   (includes duplicates)
>                  ====
> rows total       26
```

```
dim(rj2)
```

```
[1] 26  4
```

```
tail(rj2)
```

```
# A tibble: 6 x 4
   id visit Outcome  Age
  <dbl> <dbl>   <dbl> <dbl>
1     7     2    46.7  58.3
2     8     2    21.7  58.9
3     8     1     35   58.9
4     8     3    48.3  58.9
5     9    NA     NA   59.4
6    10    NA     NA    60
```

```
identical(rj2, lj) ## after some rearranging
```

```
[1] TRUE
```

Full Join

```
fj = full_join(base, visits)
```

```
Joining, by = "id"full_join: added 2 columns (visit, Outcome)
  > rows only in x      2
  > rows only in y      1
  > matched rows       24    (includes duplicates)
  >                      ====
  > rows total          27
```

```
dim(fj)
```

```
[1] 27  4
```

```
tail(fj)
```

```
# A tibble: 6 x 4
   id   Age visit Outcome
<dbl> <dbl> <dbl>   <dbl>
1     8  58.9     2    21.7
2     8  58.9     1     35
3     8  58.9     3    48.3
4     9  59.4    NA     NA
5    10   60    NA     NA
6    11   NA     3     50
```

Duplicated

- The duplicated command can give you indications if there are duplications in a **vector**:

```
duplicated(1:5)
```

```
[1] FALSE FALSE FALSE FALSE FALSE
```

```
duplicated(c(1:5, 1))
```

```
[1] FALSE FALSE FALSE FALSE FALSE TRUE
```

```
fj %>%  
  mutate(dup_id = duplicated(id))
```

```
mutate: new variable 'dup_id' with 2 unique values and 0% NA
```

```
# A tibble: 27 x 5  
   id   Age visit Outcome dup_id  
<dbl> <dbl> <dbl>   <dbl> <lgl>  
1     1   55     1     10  FALSE  
2     1   55     3    23.3  TRUE  
3     1   55     2    36.7  TRUE  
4     2  55.6     2    11.7  FALSE  
5     2  55.6     1     25   TRUE  
6     2  55.6     3    38.3  TRUE  
7     3  56.1     3    13.3  FALSE  
8     3  56.1     2    26.7  TRUE  
9     3  56.1     1     40   TRUE  
10    4  56.7     1     15  FALSE  
# ... with 17 more rows
```

Lab Part 4

[Website](#)

Finding the First (or Last) record

`pivot_longer` and `pivot_wider` are new (as of 2019) `tidyr` functions.

See link below:

<https://tidyr.tidyverse.org/dev/articles/pivot.html>

Website

[Website](#)

Reshaping data from long (tall) to wide (fat): tidyr

We can use `rowSums` to see if any values in the row is `NA` and keep if the row, which is a combination of date and line type has any non-missing data.

```
head(wide, 3)
```

```
# A tibble: 3 x 7
  day   date      daily line Alightings Average Boardings
<chr> <date>    <dbl> <chr>      <dbl>    <dbl>    <dbl>
1 Friday 2010-01-15 1644 banner      NA      NA      NA
2 Friday 2010-01-15 1644 green      NA      NA      NA
3 Friday 2010-01-15 1644 orange    1643    1644    1645
```

```
not_namat = wide %>% select(Alightings, Average, Boardings)
```

```
select: dropped 4 variables (day, date, daily, line)
```

```
not_namat = !is.na(not_namat)
head(not_namat, 2)
```

```
      Alightings Average Boardings
[1,]      FALSE      FALSE      FALSE
[2,]      FALSE      FALSE      FALSE
```

```
wide$good = rowSums(not_namat) > 0
```


Reshaping data from long (tall) to wide (fat): tidyr

Now we can filter only the good rows and delete the good column.

```
wide = wide %>% filter(good) %>% select(-good)
```

```
filter: removed 1,700 rows (37%), 2,884 rows remaining
```

```
select: dropped one variable (good)
```

```
head(wide)
```

```
# A tibble: 6 x 7
  day    date      daily line Alightings Average Boardings
  <chr> <date>    <dbl> <chr>    <dbl>    <dbl>    <dbl>
1 Friday 2010-01-15 1644 orange    1643    1644    1645
2 Friday 2010-01-22 1394. orange    1388    1394.    1401
3 Friday 2010-01-29 1332 orange    1322    1332    1342
4 Friday 2010-02-05 1218. orange    1204    1218.    1231
5 Friday 2010-02-12  671 orange     678     671     664
6 Friday 2010-02-19 1642 orange    1647    1642    1637
```

Finding the First (or Last) record

- slice allows you to select **records** (compared to first/last on a **vector**)

```
long = long %>% filter(!is.na(number) & number > 0)
```

```
filter: removed 5,364 rows (39%), 8,388 rows remaining
```

```
first_and_last = long %>% arrange(date) %>% # arrange by date
  filter(type == "Boardings") %>% # keep boardings only
  group_by(line) %>% # group by line
  slice(c(1, n())) # select ("slice") first and last (n() command) lines
```

```
filter: removed 5,630 rows (67%), 2,758 rows remaining
```

```
group_by: one grouping variable (line)
```

```
slice (grouped): removed 2,750 rows (>99%), 8 rows remaining
```

```
first_and_last %>% head(4)
```

```
# A tibble: 4 x 6
# Groups:   line [2]
  day    date      daily line  type    number
<chr>  <date>      <dbl> <chr>  <chr>    <dbl>
1 Monday 2012-06-04 13342. banner Boardings    520
2 Friday 2013-03-01     NA  banner Boardings    817
3 Tuesday 2011-11-01  8873  green  Boardings    887
4 Friday 2013-03-01     NA  green  Boardings   2592
```

Merging in base R (not covered)

Data Merging/Append in Base R

- `merge()` is the most common way to do this with data sets
 - we will use the “join” functions from `dplyr`
- `rbind/cbind` - row/column bind, respectively
 - `rbind` is the equivalent of “appending” in Stata or “setting” in SAS
 - `cbind` allows you to add columns in addition to the previous ways
- `t()` can transpose data but doesn't make it a `data.frame`

Merging

```
merged.data <- merge(base, visits, by = "id")  
head(merged.data, 5)
```

	id	Age	visit	Outcome
1	1	55.00000	1	10.00000
2	1	55.00000	3	23.33333
3	1	55.00000	2	36.66667
4	2	55.55556	2	11.66667
5	2	55.55556	1	25.00000

```
dim(merged.data)
```

```
[1] 24 4
```

Merging

```
all.data <- merge(base, visits, by = "id", all = TRUE)
tail(all.data)
```

	id	Age	visit	Outcome
22	8	58.88889	2	21.66667
23	8	58.88889	1	35.00000
24	8	58.88889	3	48.33333
25	9	59.44444	NA	NA
26	10	60.00000	NA	NA
27	11	NA	3	50.00000

```
dim(all.data)
```

```
[1] 27 4
```